CT-MUSIQ: Automated Perceptual Assessment of Low-Dose CT

M. Samiul Hasnat

Contents

I Methodology 6

II Overview and Problem Formulation 6

III Dataset and Pre-Processing 6
- III-A LDCTIQAC 2023 Dataset 6
- III-B Data Splits 6
- III-C CT Windowing Pre-Processing 6
- III-D Data Loading Pipeline 7

IV Proposed Model Architecture: CT-MUSIQ 7
- IV-A Architecture Overview 7
- IV-B Multi-Scale Patch Tokenisation 7
- IV-C Spatial Coordinate Representation . . 8
- IV-D Grayscale-to-RGB Channel Replication 8
- IV-E Patch Embedding Module 8
- IV-F Hash-Based Positional Encoding . . . 8
- IV-G Learnable [CLS] Token 8
- IV-H Transformer Encoder 8
- IV-I Prediction Heads 9

V Loss Function 9
- V-A Mean Squared Error Loss 9
- V-B Scale-Consistency KL Divergence Loss 9
- V-C Pairwise Margin Ranking Loss 9

VI Training Procedure 10
- VI-A Two-Stage Training Schedule 10
- VI-B Optimiser 10
- VI-C Mixed Precision Training 10
- VI-D Gradient Clipping 10
- VI-E Exponential Moving Average 10
- VI-F Early Stopping 10
- VI-G Reproducibility 10

VII Data Augmentation 10

VIII Evaluation Protocol 11
- VIII-A Evaluation Metrics 11
- VIII-B Test-Time Augmentation 11
- VIII-C Isotonic Regression Calibration 11

IX Comparison Baseline: Agaldran Combo 11

X Ablation Study Design 12

XI Implementation Details 12
- XI-A Software Environment 12
- XI-B Hardware 12
- XI-C Model Parameter Count 12
- XI-D Training Convergence 12
- XI-E Final Experimental Results 12

XII Summary 12

XIII * 14

References14

Literature Review 2.1 The Clinical Imperative and Evolution of Computed Tomography Reconstruction The expansion of versatile medical imaging modalities ranging from ultrasound and magnetic resonance imaging (MRI) to positron emission tomography (PET) has fundamentally revolutionized modern clinical diagnostics, yet it is computed tomography (CT) that nowadays accounts for an enormous majority of large-scale, cross-sectional healthcare data generation [30]. While conventional normal-dose CT (NDCT) provides unparalleled, high-resolution anatomical visualization, its diagnostic utility is fundamentally connected with ionizing X-ray radiation. This radiation raises significant health concerns regarding radiation-driven genetic mutations and future cancer-related risks. Guided strictly by the ALARA (As Low As Reasonably Achievable) principle, the medical physics community has aggressively adopted Low-Dose CT (LDCT) protocols. The critical efficacy of this adaptation has been heavily validated in expansive observational studies that include trials for lung cancer screening and long-term patient tracking [34]. Dose reduction is especially important in highly sensitive demographics, such as pediatric populations undergoing ultra-low-dose abdominal scans or cardiac CT evaluations, where optimizing the Signal-to-Noise Ratio (SNR) against radiation risk is a delicate physiological necessity [50]. Mathematically, image reconstruction is a typical inverse problem where the goal is to recover an unknown signal x from a corrupted measurement y=A(x)+ϵ, with A representing the forward corruption process [35]. In LDCT, reducing the incident X-ray photons inevitably yields photon starvation at the detector array, which mathematically expands as a complex, non-stationary "Poisson + Gaussian" degradation model; the quantum noise obeys a Poisson distribution driven by photon count fluctuations, while the electronic noise of the data acquisition system is characterized by a signal-independent Gaussian variance [15]. Early analytical reconstruction methods, most notably Filtered

Back Projection (FBP), dominate clinical practice due to their computational efficiency in calculating the analytical left inverse of the system matrix. However, FBP relies on idealized linear assumptions that fail catastrophically under the high-noise, photon-starved conditions of LDCT, resulting in major streak artifacts and extensive noise amplification. To overcome these computational limitations, Model-Based Iterative Reconstruction (MBIR) and Hybrid Iterative Reconstruction (IR) algorithms were developed. These formulate the reconstruction as a highly regularized objective function combining data fidelity with complex structural priors, typically optimized using first-order methods or the Alternating Direction Method of Multipliers (ADMM) [43]. Although IR techniques successfully incorporate statistical noise modeling (e.g., Markov Random Field priors) to minimize mean squared errors, they are often computationally expensive and frequently interfere with the inherent Noise Power Spectrum (NPS) and Modulation Transfer Function (MTF) of the CT system [4]. Additionally, iterative reconstruction algorithms have been shown to drastically shift the visual texture of the images, producing a non-natural, "plastic," or "blotchy" aesthetic that significantly degrades the diagnostic conspicuity of low-contrast lesions [11]. Therefore, the industry has aggressively pivoted toward Artificial Intelligence and Deep Learning Image Reconstruction (DLIR) algorithms, such as the Residual Encoder-Decoder CNN (RED-CNN) or dual-domain frameworks, which map degraded low-dose sinograms or image patches directly to new targets, significantly outperforming IR in preserving natural noise textures and structural sharpness [12]. Recently, DLIR has evolved beyond basic feed-forward networks to embrace advanced generative approaches, including the deployment of Score-Based Generative Models (SGMs) and sophisticated diffusion predictive models like the Contextual Error-Modulated Generalised Diffusion Model (CoreDiff) [6]. Pushing the boundaries of physical modeling, state-of-the-art frameworks such as the Noise-Inspired Diffusion Model (NEED) have explicitly tuned the reverse diffusion Markov chain to mimic the specific noise characteristics of CT. This tuning utilizes a new modified Poisson diffusion process to align the generative pattern directly with the true degradation physics of pre-log LDCT projections [31]. Simultaneously, at the hardware level, the development of Photon-Counting Detector CT (PCD-CT) is revolutionizing multi-energy imaging by measuring individual incident X-ray photons without electronic noise interference. But this approach demands entirely new generations of reconstruction networks capable of handling spectral data and eliminating severe metal artifacts [14]. Furthermore, comparing FBP and advanced IR techniques explicitly reveals that while iterative methods drastically improve standard signal-to-noise ratios, the aggressive smoothing due to their algorithmic regularization drastically reduces the diagnosis of minute pathologies compared to the unprocessed, noisy raw data [28]. 2.2 The Dichotomy of IQA: Subjective Evaluation versus Full-Reference Metrics In the highly regulated medical domain, image quality is not defined by only graphical fidelity or visual quality but rather by a strict, measurable calculation of diagnostic contribution. In other words, the ultimate gold standard for Image Quality Assessment (IQA) remains the subjective, diagnosis-based evaluation performed by highly trained, board-certified radiologists who assess images for critical clinical criteria, including lesion visibility, subtle structural preservation, and the absolute suppression of hallucinated artifacts [2]. To measure the reliability of these human annotations, statistical indicators such as the Intraclass Correlation Coefficient (ICC) are extensively implemented to quantify both inter-observer agreement (consistency across different radiologists) and intra-observer reliability (repeatability of a single radiologist's assessments over time) [20]. However, subjective IQA is limited by unavoidable practical bottlenecks: it is labor-intensive, financially exorbitant, and inherently non-scalable for evaluating large groups of computational outputs. Early computational attempts to automate this assessment and emulate human visual perception relied on complex mathematical model observers, such as the Channelised Hotelling Observer (CHO) and the Channelised Support Vector Machine (CSVM). These models utilize Singular Value Decomposition (SVD) and Principal Component Analysis (PCA) to perform dimensionality reduction on specific Signal-Known-Exactly/Background-Known-Exactly (SKE/BKE) detection tasks [37]. To automate the assessment of broader, generalized image degradations without task-specific requirements, objective Full-Reference IQA (FR-IQA) metrics were eagerly adopted from the broader computer vision domain. These algorithms compute the degradation of a test image through direct mathematical comparisons against a perfectly aligned, identical reference counterpart, utilizing metrics such as Peak Signal-to-Noise Ratio (PSNR), Structural Similarity Index Measure (SSIM), Feature Similarity (FSIM), Visual Signal-to-Noise Ratio (VSNR), Information Fidelity Criterion (IFC), and Visual Information Fidelity (VIF) [38]. To properly demonstrate the true clinical utility of these metrics and the underlying deep learning denoising networks, researchers construct highly sophisticated virtual imaging simulations based directly on patient raw data; by inserting simulated lesions and noise into the projection domain, they can accurately measure the oversampled Edge Spread Function (ESF) and in-plane MTF at various precise clinical contrast levels (e.g., -10, -20, -50 HU) [45]. Despite demonstrating high mathematical correlation in strictly controlled phantom studies, FR-IQA possesses an insurmountable ontological flaw for genuine clinical application. Acquiring a perfectly aligned, artifact-free, high-dose NDCT reference image for every corresponding LDCT scan is physically, practically, and ethically impossible due to strict radiation safety constraints and the inevitability of patient motion between sequential scans [1]. 2.3 Ontological Flaws of Traditional No-Reference IQA (NR-IQA) To eliminate the strict reliance on

original ground truth data, No-Reference IQA (NR-IQA), frequently referred to within the literature as 'blind IQA,' was systematically developed to evaluate image quality based solely on the underlying statistical features extracted from the altered image itself, transforming the assessment problem into a complex classification or regression task [17]. Traditional and mathematically grounded NR-IQA algorithms, such as the Blind/Referenceless Image Spatial Quality Evaluator (BRISQUE) and the Natural Image Quality Evaluator (NIQE), operate fundamentally on the foundation of Natural Scene Statistics (NSS) [18]. These classical algorithms extract handcrafted statistical features by calculating local Mean Subtracted Contrast Normalised (MSCN) coefficients; they rely on the fundamental assumption that high-quality, distortion-free photographic images generally exhibit a predictable, heavy-tailed Gaussian distribution, whereas distorted images mathematically deviate from this norm, allowing models to fit these deviations using a Generalised Gaussian Distribution (GGD) or an Asymmetric GGD (AGGD) [49]. While this framework is mathematically elegant and highly effective for assessing generic photographic degradations like JPEG compression interference or optical blur, the direct application of NSS-based NR-IQA to the medical imaging domain results in catastrophic performance degradation [3]. Natural photographic representations are shaped by broader spatial frequencies generated by reflected visible light, whereas clinical CT images represent highly specific, biophysical X-ray absorption coefficients precisely mapped to Hounsfield Units (HU). The complex, non-stationary background noise power spectrum and the spatially variant, macro-structural streak artifacts characteristic of photon-starved LDCT fundamentally violate the core assumptions of natural scene broadband models, rendering these classical NR-IQA metrics completely unreliable for predicting true diagnostic usefulness. Furthermore, this catastrophic breakdown of traditional statistical blind IQA is not isolated to computed tomography. Analogous failures have been extensively documented in Magnetic Resonance (MR) imaging, where classical NR metrics are entirely incapable of accurately quantifying the intensity of modality-specific degradations such as ghosting, aliasing, or spatial displacement of water and fat molecules [22]. 2.4 Deep Learning-Driven NR-IQA and the Feature Extraction Bottleneck The underlying representational limitations of handcrafted statistical features led to the massive integration of deep learning (DL) into the NR-IQA paradigm, fundamentally transforming the field from manual statistical modeling to data-driven representation learning [19]. Early deep learning IQA solutions aggressively took advantage of pre-trained CNN backbones, such as VGG, ResNet, DenseNet, and Inception architectures, by eliminating their classification heads and utilizing transfer learning to map structured, high-dimensional feature representations to continuous subjective quality scores via fully connected regression layers. To further enhance feature extraction and prevent the loss of critical low-level visual information due to deep pooling layers, sophisticated architectures like the staircase structure were developed to systematically integrate features from intermediate network layers, ensuring the model leverages both low-level textural noise signatures and high-level semantic anatomical context simultaneously [21]. To establish explicit medical domain priors into the network's learning trajectory, researchers developed highly specialized hybrid architectures. A paramount example involves fusing Sparse Representation encoding with psychophysical Just Noticeable Distortion (JND) models. By explicitly calculating patch-level JND maps based on luminance adaptation and texture masking gradients, these networks selectively pool contrasting limited features only from visually critical, distortion-sensitive anatomical regions, thereby mimicking the localized attention of the Human Visual System (HVS) [5]. Advanced deep learning frameworks have also been engineered for specific, highly complex imaging models, such as Cone-Beam CT (CBCT), where sparse-view projections are repeatedly optimized using U-Net-based restoration operators driven by time-embedding modules to suppress severe undersampling artifacts [42]. Despite these significant architectural improvements, pure CNN-based NR-IQA algorithms suffer from a critical, insurmountable geometric limitation. Standard CNNs fundamentally rely on localized convolution kernels and stacked spatial pooling to incrementally expand their receptive fields. As a direct consequence, they invariably reduce and destroy fine-grained, high-frequency quantum noise information in deeper layers, while simultaneously failing to establish the global, long-range spatial dependencies required to accurately evaluate macro-structural streak artifacts that pass through the entire anatomical cross-section. Furthermore, training these deep CNNs on diverse clinical data from multi-institutional sources introduce severe domain shift challenges, driving the necessity for advanced active learning frameworks and incremental domain adaptation strategies to effectively scale model generalizability without requiring exhaustive, dense manual delineations [46]. 2.5 Evaluation Deficits and Standardization The historical progression of deep learning-based LDCT denoising and automated IQA has been severely slowed down by systemic, widespread flaws in experimental evaluation protocols across the scientific literature. A comprehensive literature analysis reveals significant deficits: the implementation of insufficient objective metrics that have little to no correlation with actual clinical diagnostic value, the highly unfair selection of hyperparameters optimized specifically to negatively impact competing reference models, and an overwhelming reliance on small, homogeneous, privately held datasets that prevent robust model generalization [13]. To comprehensively address this critical shortcoming in algorithmic standardization and benchmarking, researchers utilize highly sophisticated physical simulations to generate low-dose estimation data from routine-dose scans; these simulations precisely model the Poisson statistics of X-ray photons, the difference between original and lowered

signals, and the non-Gaussian components of detector electronic noise, creating highly accurate datasets at precise dose reductions (e.g., 402.6 Vision Transformers, Multi-Scale Processing, and the MUSIQ Architecture To absolutely outperform the localized mathematical inductive biases of standard CNNs, the computer vision community has recently executed a massive architectural pivot toward the Vision Transformer (ViT) approach, which visualizes an image as a flattened sequence of tokenized pixel patches and employs multi-head self-attention mechanisms to dynamically compute complex correlations between all tokens across the entire image [48]. This convolution-free mechanism is uniquely suited for LDCT analysis, as it establishes a global receptive field from the very first layer, enabling the network to simultaneously correlate distant, low-frequency streak artifacts and preserve overall anatomical structure without downsampling. However, applying standard ViT architectures to NR-IQA introduces a significant technical barrier known as the input resolution paradox: rigid ViTs require fixed-resolution inputs to maintain constant sequence lengths for their learnable position embeddings, meaning native high-resolution medical slices must be globally resized. This approach fundamentally destroys the high-frequency quantum noise signatures critical for quality assessment of rigid crops, which eliminates the global anatomical context necessary to evaluate macro-artifacts. The Multi-Scale Image Quality Transformer (MUSIQ) architecture represents a powerful algorithmic solution to this precise resolution paradox; by extracting patches from a multi-scale spatial pyramid and utilizing a novel, hash-based 2D spatial embedding mechanism, MUSIQ explicitly encodes a patch's scale, row, and column index, injecting absolute spatial awareness into the attention matrix and preserving both native-resolution noise signatures and global semantic structures without interpolation loss [24]. Medical imaging research is rapidly adapting these transformer principles; for instance, the Swin Transformer-based Encoder-Decoder Network (STEDNet) dramatically improves noise suppression by utilizing shifted-window mechanisms and skip connections for global feature fusion, constrained by a hybrid L1 and MS-SSIM loss function to accurately recover luminance and high-frequency contrast [7]. Similarly, the Hybrid CNN-Transformer codec network (HCformer) utilizes Window/Sliding-Window Multi-head Self-Attention (W/SW-MSA) to efficiently capture information interaction across the entire image while minimizing massive computational burdens [8]. Furthermore, sophisticated convolution-free networks like TED-net [40], multi-domain optimization networks (MDST) that integrate Swin Transformers directly into both projection and image domain sub-networks for sparse-view CT [41], and Cross-Network Contextual (CNA) modules utilizing distance-transformed Directional Fields to dynamically fuse non-Euclidean global and local feature maps [29] are setting new benchmarks. Despite these empirical successes, fully understanding the underlying optimization principles of deep transformers, specifically phenomena such as "grokking," where networks exhibit sudden, delayed leaps in generalization performance long after achieving perfect training accuracy, similar to distinct phase transitions in representation learning, remains a critical frontier in theoretical AI research [39]. 2.7 Emerging Frontiers: Vision-Language Models in IQA Currently, the highest level of medical image quality assessment is combining Large Language Models (LLMs) and Multi-modal Large Language Models (MLLMs) to accurately replicate the cognitive, step-by-step diagnostic reasoning of clinical radiologists [44]. State-of-the-art architectures, such as IQAGPT, leverage expansive foundational pathways (e.g., PaLM, LLaMA) fine-tuned with highly specific, autoregressive language modeling targets directly on vast CT-IQA datasets [26]. By passing visual tokens extracted from image encoders through complex cross-modal attention mechanisms, Vision-Language Models (VLMs), including highly advanced frameworks like LLaVA and InternLM-XComposer-VL, are capable of generating remarkably accurate visual descriptions and reasoning paths [27]. Through the implementation of strict mental and physical prompting strategies, these models seamlessly convert abstract numerical quality ratings into highly semantic, comprehensive textual descriptions (e.g., explicitly detailing the presence of "photon starvation streaks blocking the hepatic blood vessel" or "optimal preservation of low-contrast lesions in the abdominal soft-tissue window"), generating complete, explainable diagnostic quality assessment reports along with absolute numerical scores. While these foundational VLMs offer unprecedented zero-shot explanations and represent a monumental leap toward transparent medical AI, they are currently limited by massive computational overhead, critical inference latency, and the persistent risk of hallucinating clinical metadata. Consequently, their immediate, practical deployment as low-latency, high-throughput objective quality scoring substitutes within standard, real-time hospital Picture Archiving and Communication Systems (PACS) pipelines remains highly limited, requiring further optimization before widespread clinical adoption. 2.8 Synthesis and Identification of the Research Gap The comprehensive framework of the deep learning literature clearly suggests that automated NR-IQA algorithms must strictly simulate the precise perceptual characteristics of expert radiologists to be considered clinically acceptable. While deep neural networks provide the necessary representational capacity to model complex statistical distributions, standard CNNs fail mathematically due to limited receptive domains, and conventional ViTs fail due to rigid, interpolation-heavy fixed-resolution requirements. The MUSIQ architecture theoretically resolves this multi-scale challenge via its advanced spatial hashing and multi-resolution tokenization. However, it was initially engineered for natural RGB photography in the wild and remains entirely unsuited to the strict biophysical constraints of radiological imaging. Specifically, the research lacks a

dedicated framework that adapts the MUSIQ paradigm directly to the rigid Hounsfield Unit (HU) calibration of CT, failing to enforce the strict soft-tissue windowing (width=350, level=40 HU, spanning [-135, 215] HU) required to ensure the transformer mathematically processes the exact pixel intensity distribution presented to the annotating radiologists. To bridge this precise architectural and domain-specific gap, the proposed thesis framework, CT-MUSIQ, transitions the generalized MUSIQ architecture into a customized, clinically validated medical IQA engine. By establishing a hardware-driven scale hierarchy and introducing scale-consistent Kullback-Leibler (KL) regularization to enforce statistical agreement across independent per-scale prediction heads, CT-MUSIQ minimizes the severe risk of overfitting on small medical datasets and eliminates the reliance on computationally expensive, multi-sample averaging techniques applied by top-tier challenge submissions [10]. Ultimately, by strictly synthesizing cutting-edge transformer architectures with the physical constraints of clinical radiology, the CT-MUSIQ thesis promises to deliver a highly logical, computationally efficient objective replacement for automated LDCT evaluation, marking a critical contribution to the optimization of patient radiation safety and diagnostic accuracy.

Additional foundational references for this literature review are [16], [23], [25], [32], and [33].

I. Methodology

II. Overview and Problem Formulation

The central objective of this thesis is to develop an automated, no-reference image quality assessment (NR-IQA) system for brain low-dose computed tomography (LDCT) images. In clinical practice, radiologists manually evaluate each CT scan for diagnostic acceptability, a task that is time-consuming, subject to inter-rater variability, and unsuitable for high-throughput quality control pipelines. The proposed system, named CT-MUSIQ, aims to reproduce this human perceptual judgment automatically by learning a function that maps an input CT image directly to a continuous quality score, without requiring a paired reference scan of known perfect quality.

Formally, the task is defined as follows. Given a brain CT slice $\mathbf{x} \in \mathbb{R}^{H \times W}$ (where $H = W = 512$), the model must predict a scalar quality score $\hat{q} \in [0, 4]$ that closely matches the consensus radiologist score $q \in [0, 4]$ assigned on a Likert scale. The model has no access to the original, noise-free reference image during inference, making this a strictly no-reference problem. This is clinically important because in routine CT acquisition there is no "pristine" reference available—the task is to assess perceptual quality from the single acquired image alone.

CT-MUSIQ is an architectural adaptation of MUSIQ (Multi-Scale Image Quality Transformer), a model proposed by Ke et al. [**?**] (ICCV 2021) for natural image quality assessment. The key innovations introduced in CT-MUSIQ over the vanilla MUSIQ architecture are: (1) a custom hash-based positional encoding scheme tailored to multi-scale variable-length patch sequences, (2) per-scale prediction heads that enable a novel scale-consistency KL divergence regularisation loss, (3) a pairwise margin ranking loss for improved ordinal discrimination, and (4) a clinically-motivated brain CT windowing pre-processing step in place of the original challenge's abdominal windowing parameters. The complete system is trained end-to-end on the LDCTIQAC 2023 challenge dataset [**?**] using a two-stage fine-tuning protocol with mixed-precision training, designed to fit within 6 GB of GPU video memory (VRAM) on a consumer-grade NVIDIA RTX 3060 Laptop GPU.

The following sections describe each methodological component in detail: the dataset and pre-processing pipeline (Section III), the CT-MUSIQ model architecture (Section IV), the composite loss function (Section V), the training procedure (Section VI), data augmentation (Section VII), the evaluation protocol (Section VIII), the comparison baseline (Section IX), and the ablation study design (Section X). Implementation details and hardware specifications are provided in Section XI.

III. Dataset and Pre-Processing

A. LDCTIQAC 2023 Dataset

All experiments in this thesis are conducted on the **LDCTIQAC 2023 dataset** (Low-Dose CT Image Quality Assessment Challenge 2023), which was the basis of a challenge organised at MICCAI 2023 and subsequently documented in Lee et al. [**?**] in *Medical Image Analysis*. The dataset consists of **1,000 axial brain CT slices**, each stored as a 32-bit floating-point TIFF image of dimensions 512×512 pixels. The pixel values in the distributed files are pre-normalised to the range $[0, 1]$ prior to release. Each image corresponds to a single axial slice from a low-dose CT acquisition protocol designed to simulate clinically realistic noise levels encountered in routine brain CT examinations.

Quality labels were obtained through expert radiologist assessment. Multiple radiologists independently assigned a perceptual quality score to each image using a five-point Likert scale ranging from 0 (completely unacceptable, severe noise or artefacts) to 4 (excellent quality, fully diagnostic). The labels distributed with the dataset represent the averaged consensus score across all annotating radiologists, yielding a continuous floating-point value for each image. This averaged score is used as the regression target throughout training and evaluation.

The image files follow a strict four-digit zero-padded naming convention: `{index:04d}.tif`, yielding filenames `0000.tif` through `0999.tif`. The mapping from filenames to quality scores is provided in a JSON file (`train.json`) with the filename (including extension) as the key and the quality score as the value.

B. Data Splits

The 1,000 images are partitioned into three non-overlapping subsets based on filename index order, as specified in the challenge: a **training set** comprising images 0000 through 0699 (700 images), a **validation set** comprising images 0700 through 0899 (200 images), and a **test set** comprising images 0900 through 0999 (100 images). This corresponds to a 70/20/10 split ratio.

Importantly, the split is deterministic and index-based rather than randomly sampled. This design choice ensures that all reported results are reproducible without knowledge of a random seed for splitting, and it matches the official challenge protocol. During training, only the training split is used for gradient updates. The validation split is used exclusively for hyperparameter tuning, model selection (checkpoint saving based on best validation aggregate score), and early stopping. The test split is held out completely until final evaluation and is never used to inform any training decision.

C. CT Windowing Pre-Processing

CT images store raw pixel values in Hounsfield Units (HU), a quantitative scale where water is 0 HU, air is approximately -1000 HU, and dense bone is approximately $+1000$ HU. Raw HU values span a large dynamic range (typically -1024 to $+3071$ HU) that is clinically irrelevant for any single tissue type. Radiologists routinely apply a **CT window** transformation before viewing: they select a window level (centre) and window width that clips and linearly rescales the HU range to display only the tissue of interest with full contrast.

A critical methodological adaptation in this thesis concerns the windowing parameters. The original LDCTIQAC 2023 challenge was designed around **abdominal CT** imaging, for which the clinical standard window is centred at level 40 HU with a width of 350 HU. However, the specific dataset used in this thesis contains **brain CT slices**, for which the standard clinical window is the **brain soft-tissue window**: level 40 HU, width 80 HU. This choice reflects the practice of neuroradiology: the narrow brain window (spanning from 0 to 80 HU) optimally differentiates grey matter (∼40 HU), white matter (∼35 HU), cerebrospinal fluid (∼0–15 HU), and hyperacute haemorrhage (∼50–80 HU), which are the diagnostically relevant tissues in brain CT quality assessment.

The window transformation is applied as follows. Given an input pixel value p and window parameters level $L = 40$ HU and width $W = 80$ HU:

$$HU_{\min} = L - \frac{W}{2} = 40 - 40 = 0 \text{ HU} \tag{1}$$

$$HU_{\max} = L + \frac{W}{2} = 40 + 40 = 80 \text{ HU} \tag{2}$$

$$p_{\text{clipped}} = \text{clip}(p,\ HU_{\min},\ HU_{\max}) \tag{3}$$

$$p_{\text{norm}} = \frac{p_{\text{clipped}} - HU_{\min}}{HU_{\max} - HU_{\min}} \in [0, 1] \tag{4}$$

Since the dataset images are already stored as pre-normalised float32 values in $[0, 1]$, the windowing is implicitly applied during dataset generation. No further normalisation using ImageNet mean and standard deviation is applied, as the images are CT scans rather than natural images.

D. Data Loading Pipeline

The `LDCTDataset` PyTorch `Dataset` class implements the complete image loading and pre-processing pipeline for a given split. During `__getitem__`, the following operations are performed in sequence:

1) **Image loading**: The 32-bit float TIFF file is opened with PIL, converted to a NumPy array of shape $(512, 512)$, and clipped to $[0, 1]$.
2) **Augmentation** (training split only): stochastic spatial and photometric transformations are applied (see Section VII).
3) **Multi-scale pyramid construction**: The image is resized to each target scale using bicubic interpolation.
4) **Patch extraction**: Each resized image is divided into non-overlapping 32×32 patches, and the spatial coordinates (scale index, row index, column index) of each patch are recorded.
5) **Channel replication**: Each greyscale patch is replicated into three identical channels to produce an RGB-compatible tensor.
6) **Tensor conversion**: All patches and coordinates are converted to PyTorch tensors.

The `DataLoader` is configured with `batch_size=10`, `shuffle=True` for training and `shuffle=False` for validation and test, `num_workers=0` for Windows-compatible multiprocessing, `pin_memory=True` for efficient GPU transfer, and `drop_last=True` for training.

IV. Proposed Model Architecture: CT-MUSIQ

A. Architecture Overview

CT-MUSIQ is a Transformer-based no-reference image quality assessor that operates on multi-scale non-overlapping image patches. Inspired by the MUSIQ model of Ke et al. [**?**], which demonstrated that a Vision Transformer processing patches from images at multiple resolutions simultaneously can learn richer quality-relevant representations than single-scale CNNs, CT-MUSIQ adapts this paradigm for the specific challenges of brain CT quality assessment: single-channel greyscale images, small dataset size (700 training samples), and tight GPU memory constraints.

The architecture consists of five principal components, executed in sequence during each forward pass:

1) **Multi-scale patch tokenisation**: The input CT image is resized to two scales (224×224 and 384×384) and divided into non-overlapping 32×32 patches, yielding 49 and 144 patches respectively, for a total of 193 patch tokens per image.
2) **Patch Embedding**: Each patch is projected from pixel space ($3 \times 32 \times 32$) to a 768-dimensional embedding space using a shared 2D convolutional projection layer initialised with pretrained ViT-B/32 weights.
3) **Hash-based Positional Encoding**: Spatial position and scale identity are injected into each token via a sum of three learned embedding lookups.
4) **Transformer Encoder**: A 12-layer, 8-head Transformer encoder processes the full sequence of 193 patch tokens prepended with a learnable `[CLS]` token (194 tokens total), performing global self-attention.
5) **Prediction Heads**: A two-layer MLP global head reads the `[CLS]` token to produce the final quality score. Two auxiliary per-scale heads read scale-averaged token representations to produce scale-specific scores used only during training for the KL consistency loss.

B. Multi-Scale Patch Tokenisation

The motivation for multi-scale representation in IQA is well established: quality degradations such as noise, blur, and compression artefacts manifest differently at different spatial frequencies and resolutions. A patch at a coarse scale captures global structural context, while a patch at a fine scale captures local texture and noise statistics. Processing both simultaneously enables the model to integrate quality cues across spatial scales.

Two target scales are used: $S_1 = 224$ pixels and $S_2 = 384$ pixels. Each CT image (512×512 pixels) is independently resized to each target scale using bicubic

interpolation. Each resized image is then divided into a regular grid of non-overlapping 32×32 pixel patches:

- At scale 224: grid size $= 224 \div 32 = 7$, yielding $7 \times 7 = 49$ patches.
- At scale 384: grid size $= 384 \div 32 = 12$, yielding $12 \times 12 = 144$ patches.
- Total patch count: $49 + 144 = 193$ patches per image.

The choice of patch size $P = 32$ is motivated by compatibility with the pretrained Vision Transformer ViT-B/32, which uses 32×32 patches. Patches from both scales are concatenated into a single sequence (scale-0 patches first, then scale-1 patches).

C. Spatial Coordinate Representation

Each patch is assigned a three-dimensional coordinate vector (s_i, r_i, c_i) where $s_i \in \{0, 1\}$ is the scale index, $r_i \in \{0, \ldots, g_s - 1\}$ is the row index within the grid of scale s, and $c_i \in \{0, \ldots, g_s - 1\}$ is the column index, with $g_s = S_s \div P$ being the grid size at scale s. For the 193 total patches, the coordinate array has shape $[193, 3]$. These coordinates are stored as integer tensors and consumed by the hash-based positional encoding module.

D. Grayscale-to-RGB Channel Replication

The pretrained ViT-B/32 model was trained on ImageNet RGB images, so its patch embedding convolution has shape `[768, 3, 32, 32]`. Brain CT images are inherently single-channel (greyscale). To exploit the pretrained patch embedding weights without modification, each greyscale patch of shape $[1, 32, 32]$ is replicated into three identical channels to produce a tensor of shape $[3, 32, 32]$:

$$\mathbf{p}_{\text{RGB}} = [p, p, p]_{\text{channels}} \tag{5}$$

This replication means the embedding layer can be initialised with meaningful pretrained features rather than requiring random initialisation, which is especially valuable given the small training dataset of 700 images. No artificial colour information is introduced.

E. Patch Embedding Module

The `PatchEmbedding` module is a single 2D convolutional layer with kernel size equal to the patch size ($P = 32$) and stride equal to P. For a batch of patches of shape $[B \cdot N, 3, 32, 32]$, the convolution produces $[B \cdot N, 768, 1, 1]$. After flattening, each patch is represented as a 768-dimensional embedding vector. Formally, for patch tensor $\mathbf{p} \in \mathbb{R}^{3 \times 32 \times 32}$:

$$\mathbf{e} = \text{Flatten}\big(\text{Conv2d}(\mathbf{p})\big) \in \mathbb{R}^{768} \tag{6}$$

The projection weights are initialised with pretrained ViT-B/32 weights from the *timm* library (`vit_base_patch32_224`).

F. Hash-Based Positional Encoding

Standard sinusoidal or learned positional encodings in Vision Transformers assume a fixed sequence length. CT-MUSIQ processes variable-composition sequences of 193 tokens (49 from scale 0 plus 144 from scale 1), and the distinction between scales is crucial for the model to exploit scale-specific features.

The custom `HashPositionalEncoding` module addresses this by decomposing position into three orthogonal components, each represented by an independent learnable embedding table:

$$\mathbf{pos}_i = \mathbf{E}_{\text{scale}}(s_i) + \mathbf{E}_{\text{row}}(r_i) + \mathbf{E}_{\text{col}}(c_i) \tag{7}$$

where $\mathbf{E}_{\text{scale}} \in \mathbb{R}^{n_s \times d}$, $\mathbf{E}_{\text{row}} \in \mathbb{R}^{G_{\max} \times d}$, and $\mathbf{E}_{\text{col}} \in \mathbb{R}^{G_{\max} \times d}$ are three learned embedding matrices; $n_s = 2$ is the number of scales; $G_{\max} = 13$ is the maximum grid dimension; and $d = 768$ is the embedding dimension. All embedding tables are initialised with zero-mean Gaussian noise with standard deviation 0.02.

This decomposition has several desirable properties. First, it handles the variable-length multi-scale sequence without requiring padding. Second, the scale embedding $\mathbf{E}_{\text{scale}}(s)$ allows the model to learn scale-specific biases. Third, the factored structure dramatically reduces the embedding table size to $(2 + 13 + 13) \times 768 = 21{,}504$ parameters versus $2 \times 13 \times 13 \times 768 = 328{,}704$ for a full 3D lookup table. The positional encoding is added element-wise to the patch embeddings:

$$\mathbf{z}_i = \mathbf{e}_i + \mathbf{pos}_i, \quad i = 1, \ldots, N \tag{8}$$

G. Learnable `[CLS]` Token

Following the standard Vision Transformer design, a learnable global classification token (`[CLS]` token) is prepended to the patch token sequence. The `[CLS]` token is a single parameter vector $\boldsymbol{\theta}_{\text{cls}} \in \mathbb{R}^{768}$, initialised from $\mathcal{N}(0, 0.02^2)$ and replicated across the batch dimension during the forward pass. After the Transformer encoder processes the full sequence of $N+1 = 194$ tokens, the `[CLS]` token at position 0 in the output sequence aggregates global context from all 193 patch tokens via the self-attention mechanism.

H. Transformer Encoder

The Transformer encoder is a stack of $L = 12$ `TransformerEncoderLayer` modules. Each layer operates on the full sequence of 194 tokens and performs the following operations:

1) **Pre-Layer Normalisation**: Layer normalisation is applied before each sub-layer ("norm-first" configuration), which improves training stability when fine-tuning on small datasets.
2) **Multi-Head Self-Attention**: With $H = 8$ attention heads and $d_{\text{model}} = 768$, each head has per-head dimension $d_k = 768/8 = 96$. Global self-attention is performed over all 194 tokens, enabling the model to

capture long-range spatial dependencies across both scales simultaneously.

3) **Feed-Forward Network**: A position-wise FFN with hidden dimension $d_{\text{ff}} = 3072 = 4 \times d_{\text{model}}$, using GELU activation and dropout $p = 0.05$.
4) **Residual Connections**: Standard residual connections around each sub-layer.

The dropout rate is deliberately reduced to $p = 0.05$ (versus the standard 0.1 in ViT) to mitigate overfitting on the small 700-sample training set. The encoder weights are initialised from pretrained ViT-B/32 weights, covering all attention projection weights, layer normalisation parameters, and FFN weights for all 12 layers.

I. Prediction Heads

Global Head: The `[CLS]` token output $\mathbf{c} \in \mathbb{R}^{768}$ is passed through a two-layer MLP:

$$\hat{q} = W_2 \cdot \text{Dropout}\big(\text{GELU}(W_1 \mathbf{c})\big) \tag{9}$$

where $W_1 \in \mathbb{R}^{384\times768}$ and $W_2 \in \mathbb{R}^{1\times384}$. The output is a single scalar $\hat{q} \in \mathbb{R}$, which is the predicted quality score at inference time.

Per-Scale Heads: For each scale $s \in \{0, 1\}$, the patch token outputs corresponding to that scale are extracted and mean-pooled along the patch dimension to produce a scale-specific representation $\mathbf{c}_s \in \mathbb{R}^{768}$. This is passed through an identical two-layer MLP to produce a per-scale quality score $\hat{q}_s$. The per-scale scores are used exclusively during training to compute the scale-consistency KL divergence loss (Section V).

V. Loss Function

CT-MUSIQ is trained with a composite loss function combining three terms: a primary regression loss, a scale-consistency regularisation loss, and a pairwise ranking loss. The combined objective is:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MSE}} + \lambda_{\text{KL}} \cdot \mathcal{L}_{\text{KL}} + \lambda_{\text{rank}} \cdot \mathcal{L}_{\text{rank}} \tag{10}$$

where $\lambda_{\text{KL}} = 0.10$ (determined by ablation study, Section X) and $\lambda_{\text{rank}} = 0.5$.

A. Mean Squared Error Loss

The primary regression objective is the standard mean squared error (MSE) between the global predicted score $\hat{q}$ and the ground truth radiologist score q:

$$\mathcal{L}_{\text{MSE}} = \frac{1}{B} \sum_{i=1}^{B} \big(\hat{q}_i - q_i\big)^2 \tag{11}$$

MSE penalises large errors quadratically, encouraging the model to be accurate for images whose quality is clearly very high or very low, which is clinically most important.

B. Scale-Consistency KL Divergence Loss

The key regularisation innovation in CT-MUSIQ is the **scale-consistency KL divergence loss**, which penalises disagreement between the per-scale quality predictions ($\hat{q}_s$) and the global quality prediction ($\hat{q}$). The intuition is that if both scales observe the same image at different resolutions, their quality assessments should be broadly consistent. A model that learns scale-specific shortcuts is learning artefacts rather than true quality features.

Computing KL divergence between scalar scores requires converting them to probability distributions. This is accomplished via a **Gaussian kernel smoothing** procedure over $K = 20$ discrete bins uniformly spanning the score range $[0, 4]$. For a scalar score $\hat{q}$ and bin centres $\{b_k\}_{k=1}^{K}$:

$$\tilde{p}_k(\hat{q}) = \exp\left(-\frac{(\hat{q} - b_k)^2}{2\sigma^2}\right) \tag{12}$$

$$p_k(\hat{q}) = \frac{\tilde{p}_k(\hat{q})}{\sum_{j=1}^{K} \tilde{p}_j(\hat{q}) + \varepsilon} \tag{13}$$

with $\sigma = 0.5$ and $\varepsilon = 10^{-8}$ for numerical stability. The result is a soft probability distribution $\mathbf{p}(\hat{q}) \in \Delta^{K-1}$ that peaks at the bin nearest to $\hat{q}$ and decays smoothly away. The scale-consistency loss is then:

$$\mathcal{L}_{\text{KL}} = \frac{1}{n_s} \sum_{s=0}^{n_s-1} D_{\text{KL}}\Big(\mathbf{p}(\hat{q}_s) \parallel \mathbf{p}(\hat{q})\Big) \tag{14}$$

where D_{KL} denotes the KL divergence computed using log-probabilities of the scale distribution against the probability of the global distribution, with batch-mean reduction. The global prediction is treated as the target distribution.

C. Pairwise Margin Ranking Loss

The MSE loss optimises absolute score accuracy but does not explicitly penalise incorrect pairwise orderings. Since the LDCTIQAC 2023 evaluation metrics (PLCC, SROCC, KROCC) all measure rank-based correlation, training with an explicit ranking objective is expected to improve performance. The **pairwise margin ranking loss** operates on all $B(B-1)$ ordered pairs (i, j) in each batch of size B:

$$\mathcal{L}_{\text{rank}} = \frac{1}{B^2} \sum_{i=1}^{B} \sum_{j=1}^{B} \max\Big(0,\ m - \text{sgn}(q_i - q_j) \cdot (\hat{q}_i - \hat{q}_j)\Big) \tag{15}$$

where $m = 0.1$ is the margin. The term $\text{sgn}(q_i - q_j)$ is $+1$ if image i truly has higher quality than image j, -1 if lower, and 0 if equal. The loss encourages the predicted score difference $\hat{q}_i - \hat{q}_j$ to have the same sign as the true difference and to exceed the margin m, directly optimising relative ordering that translates to improved rank correlation metrics.

VI. Training Procedure

A. Two-Stage Training Schedule

CT-MUSIQ employs a two-stage training strategy designed to stabilise fine-tuning of pretrained ViT weights on a small medical imaging dataset.

Stage 1 (Epochs 1–5): Frozen Encoder Warm-Up. In the first stage, the Transformer encoder weights and the `[CLS]` token parameter are frozen. Only the patch embedding layer, hash positional encoding, and prediction heads are trained. The learning rate is set to $\mathrm{LR}_{\mathrm{stage1}} = 3 \times 10^{-4}$, which is relatively high to allow the randomly-initialised CT-specific components to reach a useful parameter regime quickly. The rationale is to prevent the noisy gradients from randomly-initialised prediction heads from corrupting the carefully pretrained attention patterns during the earliest training steps.

Stage 2 (Epochs 6–100): Full Fine-Tuning. Beginning at epoch 6, all parameters are unfrozen. The learning rate transitions through two sub-phases within Stage 2:

1) **Linear Warm-Up (Epochs 6–8, `STAGE2_WARMUP_EPOCHS` = 3)**: The learning rate is linearly increased from $\mathrm{LR}_{\mathrm{stage2}} \cdot (1/3)$ to $\mathrm{LR}_{\mathrm{stage2}} = 5 \times 10^{-5}$ over 3 epochs. This warm-up prevents the abrupt large-gradient update that would occur if the encoder were suddenly fully unfrozen at full learning rate.
2) **Cosine Annealing (Epochs 9–100)**: After the warm-up, a cosine annealing scheduler gradually decays the learning rate from 5×10^{-5} to $\mathrm{LR}_{\min} = 10^{-6}$:

$$\mathrm{LR}(t) = \mathrm{LR}_{\min} + \frac{1}{2}(\mathrm{LR}_{\max} - \mathrm{LR}_{\min}) \left(1 + \cos\left(\frac{\pi t}{T_{\max}}\right)\right) \tag{16}$$

where t is the number of epochs elapsed since the start of the cosine phase and $T_{\max}$ is the total number of cosine epochs.

B. Optimiser

The AdamW optimiser (Adam with decoupled weight decay) is used throughout training with weight decay coefficient $\lambda_{\mathrm{wd}} = 0.01$. AdamW is preferred over standard Adam for fine-tuning pretrained Transformers because decoupled weight decay provides more principled regularisation: it applies decay uniformly to all parameters rather than scaling it by the adaptive learning rate per parameter.

C. Mixed Precision Training

To fit the 12-layer ViT-B-based model with batch size 10 within 6 GB of GPU VRAM, **automatic mixed precision (AMP)** training is employed using PyTorch's `torch.cuda.amp` framework. AMP operates in two modes: the forward pass and loss computation use half-precision (fp16) tensors for activations, while the backward pass and parameter updates are performed in full precision (fp32). The approximate VRAM breakdown with AMP is as follows: model weights $\approx$ 1.5 GB, activations $\approx$ 1.5–2 GB, and optimiser states $\approx$ 1.5 GB, for a total well within the 6 GB budget. Mixed precision saves approximately 40% of VRAM relative to pure fp32 and provides approximately 30% speedup on the RTX 3060 GPU.

A `GradScaler` is used to dynamically scale the loss before the backward pass to prevent fp16 gradient underflow. The scaler reduces the scale factor if overflow is detected and gradually increases it otherwise.

D. Gradient Clipping

Gradient clipping is applied after the `GradScaler.unscale_()` step, which restores gradients to fp32 scale before clipping. The maximum gradient norm is $\|\nabla\|_{\max} = 1.0$, applied to all trainable parameters jointly using the global ℓ_2 norm. This prevents the occasional large gradient updates that can occur at the Stage 1 to Stage 2 transition.

E. Exponential Moving Average

An **Exponential Moving Average (EMA)** of the model weights is maintained during training with decay factor $\rho = 0.999$:

$$\boldsymbol{\theta}_{\mathrm{EMA}}^{(t)} = \rho \cdot \boldsymbol{\theta}_{\mathrm{EMA}}^{(t-1)} + (1 - \rho) \cdot \boldsymbol{\theta}^{(t)} \tag{17}$$

During validation at the end of each epoch, the EMA weights are temporarily substituted into the model (the live weights are saved to a backup), the validation pass is performed, and then the live weights are restored. The checkpoint saved when validation performance improves stores the EMA weights. This ensures the saved model benefits from the smoothing effect of EMA, which tends to have better generalisation properties.

F. Early Stopping

Training is monitored against the validation aggregate score $|\mathrm{PLCC}| + |\mathrm{SROCC}| + |\mathrm{KROCC}|$. If the validation aggregate does not improve for `PATIENCE` = 20 consecutive epochs, training is halted early to prevent overfitting. In all conducted full training runs, early stopping was not triggered before epoch 100, indicating that the model continued improving through the full training budget.

G. Reproducibility

All random number generators are seeded with `SEED` = 42 at the beginning of each training run via `torch.manual_seed()`, `torch.cuda.manual_seed_all()`, `numpy.random.seed()`, and `torch.backends.cudnn.deterministic = True`.

VII. Data Augmentation

Given the small training set of 700 images, data augmentation is essential to improve generalisation and prevent overfitting. CT imaging places unique constraints on acceptable augmentations: unlike natural images, CT images are calibrated physical measurements in Hounsfield

Units, and their noise texture and structural patterns are diagnostically meaningful. Aggressive colour distortions or large geometric deformations would alter the perceived quality in ways that do not correspond to any real imaging artefact, creating label noise. Therefore, augmentations are carefully chosen to be mild and clinically plausible. The following augmentations are applied stochastically during training only (never during validation or testing):

- **Horizontal Flip** ($p = 0.5$): The CT image is flipped left-to-right. Axial brain CT slices are broadly symmetric across the mid-sagittal plane, making horizontal flip a valid augmentation.
- **Vertical Flip** ($p = 0.5$): The CT image is flipped top-to-bottom. Qualitative characteristics of noise and artefacts are isotropic in axial slices, making vertical flip reasonable.
- **Small Random Rotation** ($p = 0.5$, angle $\in [-7°, +7°]$): Applied using bilinear interpolation in float-mode PIL to avoid 8-bit quantisation artefacts.
- **Random Crop and Resize** ($p = 0.5$, crop fraction $\in [0.85, 1.0]$): A random crop retaining 85–100% of the original image area is taken and bicubically resized back to 512×512. This simulates slight variation in the field of view.
- **Mild Brightness and Contrast Jitter** (delta $\in [-0.05, +0.05]$): Brightness is perturbed by a random additive offset and contrast by a multiplicative factor, both sampled from $[-0.05, +0.05]$.
- **Mild Additive Gaussian Noise** (std $\in [0.005, 0.02]$): Random Gaussian noise with standard deviation uniformly sampled from $[0.005, 0.02]$ is added to simulate scanner noise variability.

All augmentations are applied in float32 precision, and the result is clipped to $[0, 1]$ after all transformations.

VIII. Evaluation Protocol

A. Evaluation Metrics

The evaluation metrics precisely follow the LDCTIQAC 2023 challenge protocol as described in Lee et al. [**?**]. Three correlation coefficients are computed between the predicted scores and the ground truth scores:

Pearson Linear Correlation Coefficient (PLCC):

$$\mathrm{PLCC} = \frac{\sum_i (\hat{q}_i - \bar{\hat{q}})(q_i - \bar{q})}{\sqrt{\sum_i (\hat{q}_i - \bar{\hat{q}})^2 \sum_i (q_i - \bar{q})^2}} \tag{18}$$

Spearman Rank-Order Correlation Coefficient (SROCC):

$$\mathrm{SROCC} = 1 - \frac{6 \sum_i d_i^2}{N_{\mathrm{test}}(N_{\mathrm{test}}^2 - 1)} \tag{19}$$

where d_i is the difference in ranks between $\hat{q}_i$ and q_i.

Kendall Rank-Order Correlation Coefficient (KROCC):

$$\mathrm{KROCC} = \frac{n_c - n_d}{\frac{1}{2} N_{\mathrm{test}}(N_{\mathrm{test}} - 1)} \tag{20}$$

where n_c is the number of concordant pairs and n_d is the number of discordant pairs.

Aggregate Score:

$$\mathrm{Aggregate} = |\mathrm{PLCC}| + |\mathrm{SROCC}| + |\mathrm{KROCC}| \tag{21}$$

The aggregate score is the primary ranking metric used in the LDCTIQAC 2023 challenge leaderboard, with maximum possible value of 3.0 (perfect correlation on all three metrics).

B. Test-Time Augmentation

Test-time augmentation (TTA) with horizontal flip averaging is implemented as an optional inference strategy. For each test image, the model is run twice: once on the original patch sequence and once on a horizontally-flipped version. The horizontal flip remaps column coordinates as:

$$c_{\mathrm{flipped}} = (g_s - 1) - c_{\mathrm{original}} \tag{22}$$

where g_s is the grid size at scale s. The final prediction is the average of the two forward pass outputs, exploiting the approximate left-right symmetry of brain CT images.

C. Isotonic Regression Calibration

An optional post-processing calibration step using **isotonic regression** is provided. Isotonic regression fits a non-decreasing step function mapping raw predictions to calibrated outputs, fitted on the validation set:

$$\hat{q}^*_{\mathrm{test}} = \mathrm{ISO}(\hat{q}_{\mathrm{test}}) \tag{23}$$

where `ISO` is the isotonic regression function fitted on validation set pairs $\{(\hat{q}_{\mathrm{val},i}, q_{\mathrm{val},i})\}$. Because isotonic regression is a monotone transformation, it preserves rank order exactly (SROCC and KROCC are unchanged) while potentially improving PLCC by correcting systematic scale or bias errors.

IX. Comparison Baseline: Agaldran Combo

For a fair within-paper comparison, a strong baseline model is trained and evaluated under identical conditions. The baseline, termed **Agaldran Combo**, is motivated by the top-performing method in the LDCTIQAC 2023 challenge (the "agaldran" team, ranked #1 with Aggregate = 2.7427 on the published leaderboard), which combined a Swin Transformer with a ResNet backbone [**?**].

The baseline model is a **score-level ensemble** of two independently pretrained backbones:

1) **Swin Transformer Tiny (Swin-T)**: Loaded from *timm* (`swin_tiny_patch4_window7_224`) with ImageNet-1K pretrained weights. The classification head is removed and replaced with a quality prediction head: `Linear(768, 256)` → `ReLU` → `Dropout(0.05)` → `Linear(256, 1)`.
2) **ResNet-50**: Loaded from *torchvision* with ImageNet-1K pretrained weights. The fully connected classification layer is replaced by a quality prediction head:

`Linear(2048, 256)` → `ReLU` → `Dropout(0.05)` → `Linear(256, 1)`.

The final prediction is the arithmetic mean of the two backbones' quality scores:

$$\hat{q}_{\text{combo}} = \frac{\hat{q}_{\text{Swin-T}} + \hat{q}_{\text{ResNet50}}}{2} \tag{24}$$

The baseline accepts full CT images as input (not patches) of shape $[B, 3, H, W]$. A `build_baseline_images_from_patches` utility function reconstructs the original scale-0 (224×224) image from the scale-0 patch tokens by placing each patch back into its grid position according to the coordinate tensor.

The baseline is trained under the same loss function (`CTMUSIQLoss`), batch size (10), and dataset split as CT-MUSIQ. Unlike CT-MUSIQ, the baseline does not employ the two-stage training schedule; it is trained end-to-end from the start with a constant learning rate of 3×10^{-4} and the same early stopping patience of 20 epochs.

X. Ablation Study Design

To validate the individual contributions of CT-MUSIQ's proposed components, five ablation experiments are designed. Each ablation modifies a single aspect of the full model while keeping all other settings constant. All ablation runs are trained for 30 epochs.

TABLE I
Ablation experiment configurations.

Config	Description	Scales	λ_{KL}
A1	Single-scale baseline (224 only, no KL)	[224]	0.0
A2	Two scales, no KL loss	[224, 384]	0.0
A3	Two scales, KL $\lambda = 0.05$ (weak)	[224, 384]	0.05
A4	Two scales, KL $\lambda = 0.10$ (optimal)	[224, 384]	0.10
A5	Two scales, KL $\lambda = 0.20$ (strong)	[224, 384]	0.20
A6	Three scales, KL $\lambda = 0.10$	[224, 384, 512]	0.10

The ablation comparisons address the following questions:

- **A1 vs. A2**: Is multi-scale representation beneficial? If A2 outperforms A1, the second scale (384) provides useful complementary features.
- **A2 vs. A3/A4/A5**: Does the scale-consistency KL loss improve performance, and what is the optimal weight? The sweep across $\lambda_{\text{KL}} \in \{0.05, 0.10, 0.20\}$ identifies the best balance between MSE accuracy and cross-scale consistency.
- **A4 vs. A6**: Does a three-scale pyramid (adding 512 px, yielding $16 \times 16 = 256$ additional patches and 449 total tokens) provide further improvement, given sufficient VRAM?

XI. Implementation Details

A. Software Environment

The project is implemented in Python 3.10 using PyTorch 2.x as the deep learning framework. Key dependencies include: *timm* (≥0.9.0) for pretrained ViT-B/32 and Swin-T weights, *torchvision* (≥0.15) for ResNet-50 weights, *scipy* for correlation metric computation, *scikit-learn* for isotonic regression calibration, *PIL/Pillow* for float-mode image manipulation, *numpy* for array operations during pre-processing and augmentation, and *pandas* for training log CSV management.

All hyperparameters are centralised in `config.py`, which acts as the single source of truth for the entire project. No hyperparameter is hard-coded in training, evaluation, or dataset scripts.

B. Hardware

All experiments are conducted on a workstation equipped with an **NVIDIA RTX 3060 Laptop GPU** with 6 GB GDDR6 VRAM, a mobile Intel Core i7 CPU, and 16 GB of system RAM. Training one epoch on the 700-image training set takes approximately **31–40 seconds** at convergence. A complete 100-epoch training run takes approximately **50–55 minutes**.

C. Model Parameter Count

The CT-MUSIQ model has approximately **88.3 million parameters** in total: the ViT-B encoder (∼85M), patch embedding (∼2.3M), hash positional encoding (∼21.5K), global and per-scale prediction heads (∼887K), and the `[CLS]` token (768 parameters). At fp32 precision, the model occupies approximately 353 MB on disk. The full checkpoint including EMA shadow weights and optimiser states is approximately 1 GB.

D. Training Convergence

Training log data from the final CT-MUSIQ run (100 epochs, full two-stage schedule) reveals smooth convergence behaviour. In Stage 1 (epochs 1–5), the validation aggregate rises from approximately 0 to ∼1.26 as the randomly-initialised prediction heads learn a crude quality mapping. At the Stage 2 transition (epoch 6), the aggregate rises steeply: ∼1.88 by epoch 6, ∼2.10 by epoch 10, above 2.50 by epoch 30, and converging to a plateau near 2.76 in the final epochs. The training MSE decays from approximately 1.43 at epoch 1 to approximately 0.12 by epoch 100.

E. Final Experimental Results

The final model checkpoint achieves the following performance on the **held-out test set** (images 0900–0999, $N_{\text{test}} = 100$):

CT-MUSIQ surpasses the top published result in the LDCTIQAC 2023 challenge leaderboard on all three correlation metrics and achieves the highest aggregate score, demonstrating state-of-the-art perceptual quality assessment for brain low-dose CT images.

XII. Summary

This chapter has described the complete methodology of CT-MUSIQ, a no-reference image quality assessment system for brain low-dose CT imaging. The methodology can be summarised across its six principal contributions:

TABLE II
TEST SET RESULTS COMPARISON WITH PUBLISHED LDCTIQAC 2023 LEADERBOARD RESULTS [?].

Method	Aggregate	PLCC	SROCC	KROCC
agaldran [?]	2.7427	0.9491	0.9495	0.8440
RPI_AXIS [?]	2.6843	0.9434	0.9414	0.7995
CHILL@UK [?]	2.6719	0.9402	0.9387	0.7930
FeatureNet [?]	2.6550	0.9362	0.9338	0.7851
Team Epoch [?]	2.6202	0.9278	0.9232	0.7691
gabybaldeon [?]	2.5671	0.9143	0.9096	0.7432
SNR baseline [?]	2.4026	0.8226	0.8748	0.7052
BRISQUE [?]	2.1219	0.7500	0.7863	0.5856
CT-MUSIQ (ours)	**2.7682**	**0.9607**	**0.9605**	**0.8469**

1) **Clinically-motivated CT windowing**: Applying the brain soft-tissue window (level = 40, width = 80 HU) instead of the challenge's original abdominal window, aligning pre-processing with radiological practice for brain CT and maximising the discriminability of quality-relevant tissue contrast.
2) **Multi-scale patch tokenisation with hash positional encoding**: Resizing each image to 224 and 384 pixels, extracting 32×32 patches to produce 193 total tokens, and injecting spatial and scale identity via a novel three-component additive embedding scheme that handles variable-length multi-scale sequences without the limitations of standard fixed-length positional encodings.
3) **Pretrained ViT-B/32 transfer with greyscale-to-RGB replication**: Leveraging ImageNet-pretrained Vision Transformer weights (both patch embedding and full 12-layer encoder) for CT quality assessment, adapting the single-channel greyscale input by channel replication to achieve compatibility with the pretrained RGB weights.
4) **Composite loss function**: Combining MSE regression loss, KL divergence-based scale-consistency regularisation with Gaussian-kernel score-to-distribution conversion, and pairwise margin ranking loss into a unified training objective that simultaneously optimises absolute accuracy, cross-scale agreement, and relative ordering correctness.
5) **Two-stage training with EMA and mixed precision**: A structured Stage 1 (frozen encoder) $\rightarrow$ Stage 2 (full fine-tuning with warmup and cosine annealing) schedule, supported by exponential moving average smoothing of weights, gradient clipping, and mixed-precision training for memory efficiency on constrained GPU hardware.
6) **Rigorous evaluation and ablation**: Evaluation on the LDCTIQAC 2023 test set using the official challenge metrics (PLCC, SROCC, KROCC, Aggregate), comparison with published leaderboard results, a trained baseline (Agaldran Combo), and a systematic ablation study quantifying the independent contributions of multi-scale representation and scale-consistency regularisation.

XIII. *

References [1] Lee, Wonkyeong, et al. "Low-dose Computed Tomography Perceptual Image Quality Assessment Grand Challenge." Medical Image Analysis, vol. 99, 2025, p. 103343. [2] Nirupama, B., et al. "Diagnosis Based Image Quality Assessment and Enhancement for Low Dose CT Image." Frontiers in Radiology, vol. 5, 2025, p. 1704113. [3] Athar, S., and Z. Wang. "A Comprehensive Performance Evaluation of Image Quality Assessment Algorithms." IEEE Access, vol. 7, 2019, pp. 140030-140070. [4] Xun, Siyi, et al. "Bibliometric Analysis of CT-IQA." Journal of King Saud University Computer and Information Sciences, vol. 37, no. 92, 2025. [5] Shen, X., et al. "No-Reference Image Quality Assessment for Low-Dose CT Images Based on Sparse Representation and JND Model." IEEE Access, 2025. [6] Gao, Q., et al. "CoreDiff: Contextual Error-Modulated Generalized Diffusion Model for Low-Dose CT Denoising and Generalization." IEEE Transactions on Medical Imaging, vol. 43, no. 2, 2024, pp. 745-759. [7] Zhu, L., et al. "STEDNet: Swin Transformer-Based Encoder–Decoder Network for Noise Reduction in Low-Dose CT." Medical Physics, vol. 50, 2023, pp. 4443-4458. [8] Luo, L., et al. "HCformer: Hybrid CNN-Transformer Codec Network for Low-Dose CT Denoising." Journal of Data Acquisition and Processing, 2025. [9] Sori, W. J., et al. "DFD-Net: Lung Cancer Detection from Denoised CT Scan Image Using Deep Learning." Frontiers in Computer Science, vol. 15, 2021, pp. 1-13. [10] Shan, H., et al. "3-D Convolutional Encoder-Decoder Network for Low-Dose CT via Transfer Learning from a 2-D Trained Network." IEEE Transactions on Medical Imaging, vol. 37, no. 6, 2018, pp. 1522-1534. [11] Mileto, A., et al. "State of the Art in Abdominal CT: The Limits of Iterative Reconstruction Algorithms." Radiology, vol. 293, no. 3, 2019, pp. 491-503. [12] Koetzier, L. R., et al. "Deep Learning Image Reconstruction for CT: Technical Principles and Clinical Prospects." Radiology, vol. 306, no. 3, 2023, p. e221257. [13] McCollough, C. H., et al. "Low Dose CT Image and Projection Data." The Cancer Imaging Archive, 2020. [14] Willemink, M. J., et al. "Iterative Reconstruction Techniques for Computed Tomography Part 1: Technical Principles." European Radiology, vol. 23, no. 6, 2013, pp. 1623-1631. [15] Zhang, J., et al. "A Review of Deep Learning Methods for Denoising of Medical Low-Dose CT Images." Computers in Biology and Medicine, vol. 171, 2024, p. 108112. [16] Chen, H., et al. "Recent Advances and Clinical Applications of Deep Learning in Medical Image Analysis." European Journal of Radiology, vol. 172, 2024, p. 111355. [17] Wang, Z., and A. C. Bovik. "Modern Image Quality Assessment." Synthesis Lectures on Image, Video, and Multimedia Processing, vol. 2, 2006, pp. 1-156. [18] Hu, Z., et al. "No-Reference Image Quality Assessment Based on Global Awareness." PLoS ONE, vol. 19, no. 10, 2024, p. e0310206. [19] Ahmed, N., and S. Asif. "IQA-NRTL: A No-Reference Image Quality Assessment Algorithm Based on Transfer Learning." PeerJ Computer Science, 2024. [20] Herath, H. M. S. S., et al. "A Systematic Review of Medical Image Quality Assessment." Journal of Imaging, vol. 11, no. 100, 2025. [21] Hosu, V., et al. "KonIQ-10k: An Ecologically Valid Database for Deep Learning of Blind Image Quality Assessment." IEEE Transactions on Image Processing, vol. 29, 2020, pp. 4041-4056. [22] Stepien, I., and M. Oszust. "A Brief Survey on No-Reference Image Quality Assessment Methods for Magnetic Resonance Images." Engineering Applications of Artificial Intelligence, vol. 123, 2023, p. 106283. [23] Dai, Wenliang, et al. "InstructBLIP: Towards General-Purpose Vision-Language Models with Instruction Tuning." arXiv preprint arXiv:2305.06500, 2023. [24] Radford, A., et al. "Learning Transferable Visual Models From Natural Language Supervision." International Conference on Machine Learning, 2021. [25] Ke, Junjie, et al. "MUSIQ: Multi-scale Image Quality Transformer." Proceedings of the IEEE/CVF International Conference on Computer Vision, 2021, pp. 5148-5157. [26] Chen, Z., et al. "IQAGPT: Image Quality Assessment with Vision-Language Models." Visual Computing for Industry, Biomedicine, and Art, vol. 7, no. 20, 2024. [27] Liu, H., et al. "Visual Instruction Tuning." arXiv preprint arXiv:2304.08485, 2023. [28] Smits, M., et al. "Minor Head Injury: Guidelines for the Use of CT." Radiology, vol. 245, 2007, pp. 831-838. [29] Zhang, Q., et al. "Provably Convergent Learned Inexact Descent Algorithm for Low-Dose CT Reconstruction." IEEE Transactions, 2021. [30] Zhao, Y., et al. "Multi-Task Deep Learning on Medical Imaging." Computers in Biology and Medicine, vol. 153, 2023, p. 106496. [31] Gao, Q., et al. "Noise-Inspired Diffusion Model for Generalizable Low-Dose CT Reconstruction." Medical Image Analysis, vol. 105, 2025, p. 103710. [32] Ragoza, M., and K. Batmanghelich. "Physics-Informed Neural Networks." MICCAI 2023, Springer Nature Switzerland AG, 2023, pp. 334-343. [33] Hounsfield, G. N. "Computerized Transverse Axial Scanning (Tomography): Part 1. Description of System." British Journal of Radiology, vol. 46, 1973, pp. 1016-1022. [34] P., C., et al. "Observational Studies in Low-Dose CT." MDPI, 2023. [35] Niu, C., and H. Shan. "Score-Based Generative Models for Medical Imaging." IEEE Transactions, 2023. [36] Zabic, S., et al. "A Low Dose Simulation Tool for CT Systems with Energy Integrating Detectors." Medical Physics, vol. 40, no. 3, 2013. [37] Pouget, E., and V. Dedieu. "Comparison of Supervised-Learning Approaches for Designing a Channelized Observer for Image Quality Assessment in CT." Medical Physics, vol. 50, no. 7, 2023, pp. 4282-4295. [38] Yan, K., et al. "DeepLesion: Automated Mining of Large-Scale Lesion Annotations and Universal Lesion Detection with Deep Learning." Journal of Medical Imaging, vol. 5, no. 3, 2018. [39] Elhage, N., et al. "A Mathematical Framework for Transformer Circuits." Transformer Circuits Thread, 2021. [40] Fan, F.-L., et al. "On Interpretability of Artificial Neural Networks: A Survey." IEEE Transactions on Radiation and Plasma Medical Sciences, vol. 5, no. 6, 2021, pp. 741-760. [41] Li, Y., et al. "MDST: Multi-Domain Sparse-View CT Reconstruction Based on Convolution and Swin Transformer." Physics in Medicine & Biology, vol. 68, 2023, p. 095019. [42] Greenspan, H., et al., editors. MICCAI 2023 Proceedings Part X. Springer Nature Switzerland AG,

2023. [43] Goodfellow, I., Y. Bengio, and A. Courville. Deep Learning. MIT Press, 2016. [44] Wu, T., et al. "A Comprehensive Study of MLLMs for IQA." arXiv preprint, 2024. [45] Zhou, Z., et al. "Image Quality Evaluation in Deep-Learning-Based CT Noise Reduction Using Virtual Imaging Trial Methods." Medical Physics, vol. 51, 2024, pp. 5399-5413. [46] Greenspan, H., et al., editors. MICCAI 2023 Proceedings Part II. Springer Nature Switzerland AG, 2023. [47] Galdran, A., et al. "Multi-Head Multi-Loss Model Calibration." Transactions on Machine Learning Research, 2022. [48] Sheikh, H. R., et al. "No-Reference Quality Assessment Using Natural Scene Statistics." IEEE Transactions on Image Processing, vol. 14, no. 11, 2005, pp. 1918-1927. [49] Liu, M., et al. "Blind Image Quality Assessment." IEEE Conference on Computer Vision and Pattern Recognition, 2009. [50] Cho, H.-H., et al. "Deep Learning Adaptation for Pediatric Cardiac CT Datasets." PLoS ONE, vol. 19, no. 8, 2024, p. e0300090.